

TensorFlow Implementaton of Laplace Approximation for GP Hyperparameters (w/ Internal Newton Optimiser)

This is an example python script to illustrate LA GP fitting with TensorFlow

Step 1: Read in summarised TRaC data with matched covariates and build as TensorFlow constants

- Age restricted to 1-50 y/o (to limit impact of maternal anti-bodies and limit look-back time)
- Sites snapped to land-sea mask boundary and standardised covariates extracted
- Area including Port-au-Prince and surrounds excluded (to limit impact of migration)

In [1]:

```
import numpy as np
import pandas as pd
import tensorflow as tf

tracdata = pd.read_csv('./Data/summary_TRaC_data_with_covariates.csv')

meanMSP = np.log(tracdata.groupby(['Cluster_Num'])['MSP'].mean().values) # In this contrived example we will fit a GP regression for the mean MSP in each village
Nvillages = meanMSP.size
meanMSP = tf.constant(meanMSP, dtype='float32')
covariates = tracdata.groupby(['Cluster_Num'])['covariate_accessibility', 'covariate_AI', 'covariate_distTowater', 'covariate_elevation', 'covariate_forest', 'covariate_grass', 'covariate_urbanbarren', 'covariate_woodysavanna', 'covariate_OSM', 'covariate_PET', 'covariate_slope', 'covariate_TWI'].mean().values
Ncovariates = covariates.shape[1]
covariates = tf.constant(covariates, dtype='float32')
```

Step 2: Read in INLA projection and SPDE precision matrices

- These are pre-cooked in the R preprocessing script

In [2]:

```
from scipy.io import mmread
Amatrix = mmread('./Data/A_matrix.mtx')
Nmesh = Amatrix.shape[1]
M0 = mmread('./Data/M0_matrix.mtx')
M1 = mmread('./Data/M1_matrix.mtx')
M2 = mmread('./Data/M2_matrix.mtx')
Amatrix = tf.constant(Amatrix.toarray(), dtype='float32') # Not sure if it's worth figuring out how to use SparseTensors
M0 = tf.constant(M0.toarray(), dtype='float32')
M1 = tf.constant(M1.toarray(), dtype='float32')
M2 = tf.constant(M2.toarray(), dtype='float32')
```

Step 3: Priors and transformations for GMRF model

- Precision matrix (Q) constructed as per Eqn 10 of Lingren et al.:
<https://rss.onlinelibrary.wiley.com/doi/full/10.1111/j.1467-9868.2011.00777.x>
(<https://rss.onlinelibrary.wiley.com/doi/full/10.1111/j.1467-9868.2011.00777.x>)

In [3]:

```
log_kappa = tf.Variable(1.0, dtype='float32')
prior_log_kappa = tf.contrib.distributions.Normal(loc = 0.0, scale = 1.0)

log_tau = tf.Variable(-3.0, dtype='float32')
prior_log_tau = tf.contrib.distributions.Normal(loc = 0.0, scale = 1.0)

log_field = tf.constant(np.zeros(Nmesh), dtype='float32')
spde_prec = (tf.exp(log_kappa*4.0)*M0 + tf.constant(2.0, dtype='float32')*tf.exp(log_kappa*2.0)*M1 + M2)*tf.exp(log_tau*tf.constant(2.0, dtype='float32'))
spde_cov = tf.matrix_inverse(spde_prec)
prior_log_field = tf.contrib.distributions.MultivariateNormalFullCovariance(loc = tf.constant(np.zeros(Nmesh), dtype='float32'), covariance_matrix = spde_cov)

slopes = tf.Variable(np.zeros(Ncovariates), dtype='float32')
prior_slopes = tf.contrib.distributions.Normal(loc = 0.0, scale = 1.0)

intercept = tf.Variable(0.0, dtype='float32')
prior_intercept = tf.contrib.distributions.Normal(loc = 0.0, scale = 5.0)
```

Step 4: Set up empirical Bayes GP hyperparameter search in TensorFlow

In [4]:

```
def inner_vector_update(log_field, iter_diff, log_laplace_approximation, output_collect
ions=(), name=None):
    log_field = tf.convert_to_tensor(log_field)

    linear_predictor = tf.squeeze(tf.matmul(Amatrix,tf.expand_dims(log_field,1))) + int
ercept + tf.squeeze(tf.matmul(covariates,tf.expand_dims(slopes,1)))
    log_prior = prior_log_kappa.log_prob(log_kappa) + prior_log_tau.log_prob(log_tau) +
prior_log_field.log_prob(log_field) + tf.reduce_sum(prior_slopes.log_prob(slopes)) + p
rior_intercept.log_prob(intercept)
    likelihood_data = tf.contrib.distributions.Normal(loc = linear_predictor, scale = 1
0.0)
    log_likelihood = tf.reduce_sum(likelihood_data.log_prob(meanMSP))
    negative_log_posterior_prob = -(log_likelihood + log_prior)

    new_log_field = log_field - tf.constant(0.5, dtype = 'float32')*tf.squeeze(tf.matm
l(tf.squeeze(tf.matrix_inverse(tf.hessians(negative_log_posterior_prob,log_field))),tf.
expand_dims(tf.squeeze(tf.gradients(negative_log_posterior_prob,log_field)),1)))
    new_iter_diff = tf.reduce_sum(tf.abs(log_field-new_log_field))

    new_log_laplace_approximation = tf.constant(Nmesh*0.9189385, dtype = 'float32') - t
f.constant(0.5, dtype = 'float32')*tf.linalg.logdet(tf.squeeze(tf.hessians(negative_log
_posterior_prob,log_field))) - negative_log_posterior_prob

    return [new_log_field, new_iter_diff, new_log_laplace_approximation]

def running_condition(log_field,iter_diff,log_laplace_approximation, output_collections
=(), name=None):
    return tf.greater(iter_diff,tf.constant(0.001, dtype = 'float32'))

log_laplace_approximation = tf.constant(0.0, dtype = 'float32')
iter_diff = tf.constant(1000000000000000000.0, dtype = 'float32')

iteration_outputs = tf.while_loop(running_condition,inner_vector_update,[log_field,iter
_diff,log_laplace_approximation])
```

Step 5: Run optimisation to find empirical Bayes hyperparameters

- Uses the AdaGrad algorithm: probably we want to switch to our own Bayesian optimisation later(?)

In [5]:

```
log_laplace_approx_at_conditional_mode = iteration_outputs[2]
log_field_at_conditional_mode = iteration_outputs[0]

optimiser = tf.train.AdagradOptimizer(0.05) # Behaviour sensitive to this training rate!
train = optimiser.minimize(-log_laplace_approx_at_conditional_mode, var_list = [log_kappa, log_tau, slopes, intercept])

sess = tf.Session()
sess.run(tf.global_variables_initializer())

sess.run(train)
sess.run(train)

log_kappa_diff = 1000.0 # This is not a sophisticated stopping criterion!
old_log_kappa = sess.run(log_kappa)
while (log_kappa_diff > 0.01):
    sess.run(train)
    log_kappa_diff = np.abs(sess.run(log_kappa)-old_log_kappa)
    old_log_kappa = sess.run(log_kappa)
    print(log_kappa_diff)

0.024695098
0.021266878
0.018945813
0.017234087
0.015896618
0.014811575
0.013916552
0.013155341
0.012502015
0.011930466
0.011422396
0.010971367
0.010562897
0.010193765
0.009855807
```

Step 6: Plot posterior mode field at empirical Bayes hyperparameters

- Requires GDAL for raster io: <https://www.lfd.uci.edu/~gohlke/pythonlibs/#gdal> (<https://www.lfd.uci.edu/~gohlke/pythonlibs/#gdal>)

In [6]:

```
log_field = sess.run(log_field_at_conditional_mode)
slopes = sess.run(slopes)
intercept = sess.run(intercept)

Afullmatrix = mmread('./Data/Afull_matrix.mtx')
Afullmatrix = Afullmatrix.toarray()
covariatesfull = pd.read_csv('./Data/fullcovariates.csv').values[:,1:13]
linear_predictor = np.matmul(Afullmatrix,log_field) + np.matmul(covariatesfull,slopes)
+ intercept

import gdal
import os
ds = gdal.Open('C:/Users/zool1232/Work/Code/LA_Example/CovariateRasters/covariates_AI.tif')
covariate = ds.GetRasterBand(1)
nodata = covariate.GetNoDataValue()
referenceimage = covariate.ReadAsArray()
referenceimage = np.ma.masked_less(referenceimage,-3)

valid_pixels = referenceimage.mask==0
valid_pixels_indices = np.nonzero(np.ndarray.flatten(valid_pixels)) # row-major flattening
n_valid_pixels = np.sum(valid_pixels)

xoff, a, b, yoff, d, e = ds.GetGeoTransform() # NOTE: this seems like an incredibly tedious way to deal with raster geometries, but I could not find anything more sensible on line (??)
pixel_longs = a * np.transpose(np.reshape(np.repeat(np.arange(referenceimage.shape[1]),referenceimage.shape[0]),(referenceimage.shape[1],referenceimage.shape[0]))) + b * np.reshape(np.repeat(np.arange(referenceimage.shape[0]),referenceimage.shape[1]),referenceimage.shape) + xoff
pixel_lats = d * np.transpose(np.reshape(np.repeat(np.arange(referenceimage.shape[1]),referenceimage.shape[0]),(referenceimage.shape[1],referenceimage.shape[0]))) + e * np.reshape(np.repeat(np.arange(referenceimage.shape[0]),referenceimage.shape[1]),referenceimage.shape) + yoff

referenceimage.put(np.where(np.ndarray.flatten(referenceimage.mask)==0),linear_predictor)
```

In [8]:

```
import matplotlib.pyplot as pyplot
imgplot = pyplot.imshow(referenceimage, extent=[np.min(pixel_longes[valid_pixels]), np.max(
(pixel_longes[valid_pixels]), np.min(pixel_lats[valid_pixels]), np.max(pixel_lats[valid_pi
xels]))])
#pyplot.scatter(longs, lats, c='red') # check image coordinate transforms by comparing
raster plot against observation locations
pyplot.show() # NOTE: seems to only display figure if I run this code a second time aro
und: why?
```

